

# Introduction to Deep Learning

## Chapter 1: Introduction

Pao-Ann Hsiung

National Chung Cheng University

# Contents

- ▶ Chapter Goals
- ▶ What is machine/deep learning?
- ▶ History of neural networks and deep learning
- ▶ Applications of deep learning
- ▶ Techniques required
- ▶ What we learnt?
- ▶ What to do now?

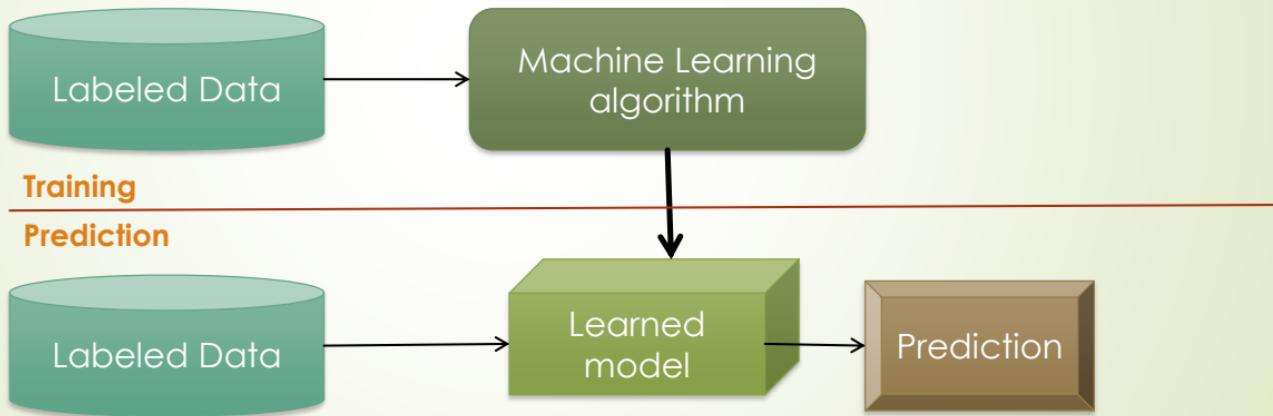
# Chapter Goals

## -- Introduction

- ▶ To be knowledgeable about what exactly **machine** and **deep learning** is
- ▶ To be familiar with the **history of deep learning**,
- ▶ To be familiar with the **applications of deep learning**,
- ▶ To be familiar with the **techniques required for deep learning**
- ▶ To learn how to program in the **Python** programming language

# What is machine learning (ML)?

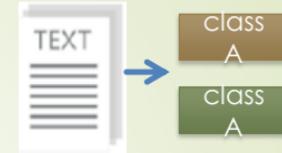
- Machine learning is a field of computer science that gives computers the ability to **learn without being explicitly programmed**



Methods that can learn from and make predictions on data

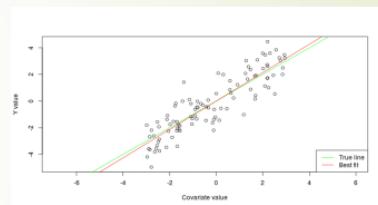
# Types of Learning

**Supervised:** Learning with a **labeled training** set of data  
 Example: learn the **classification** of images based on image **labels** (dogs/cats, day time, numbers, etc.)



Classification

**Unsupervised:** Discover **patterns** in **unlabeled** data  
 Example: **cluster** similar documents based on text



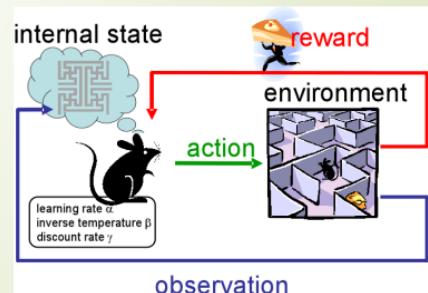
Regression



Clustering

**Reinforcement learning:** learn to **act** based on **feedback/reward**

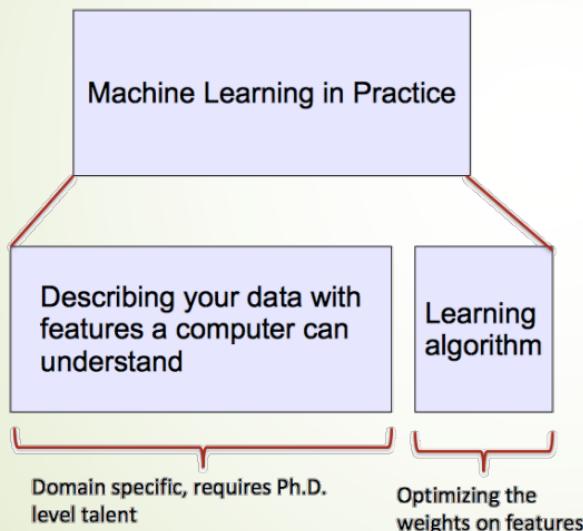
Example: learn to play Go, reward: **win or lose**



# Machine Learning in Practice

Most machine learning methods work well because of human-designed representations and input features

ML becomes just **optimizing weights** to best make a final prediction



Feature	NER
Current Word	✓
Previous Word	✓
Next Word	✓
Current Word Character n-gram	all
Current POS Tag	✓
Surrounding POS Tag Sequence	✓
Current Word Shape	✓
Surrounding Word Shape Sequence	✓
Presence of Word in Left Window	size 4
Presence of Word in Right Window	size 4

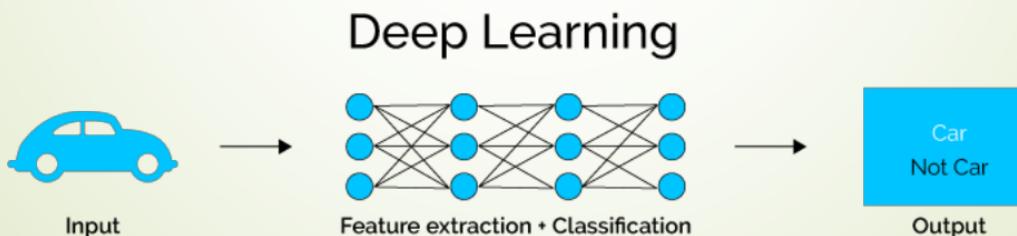
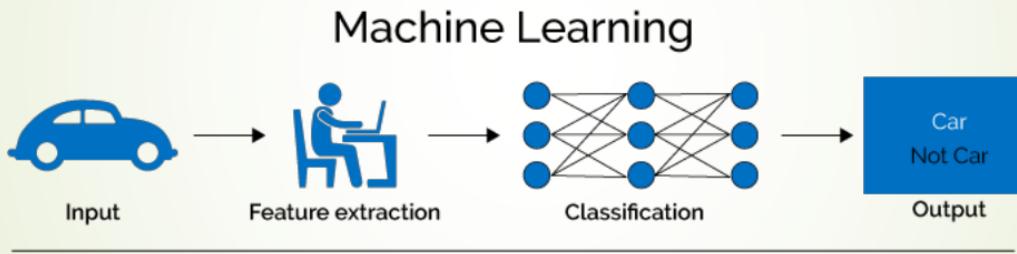
NER: Named Entity Recognition  
POS: Part of Speech

# Feature Hierarchy

- ▶ Hierarchy of representations with increasing levels of abstraction
- ▶ Image recognition
  - ▶ Pixel → edge → texton → motif → part → object
- ▶ Text
  - ▶ Character → word → word group → clause → sentence → story
- ▶ Speech
  - ▶ Sample → spectral band → sound → ... → phone → phoneme → word

# What is Deep Learning (DL) ?

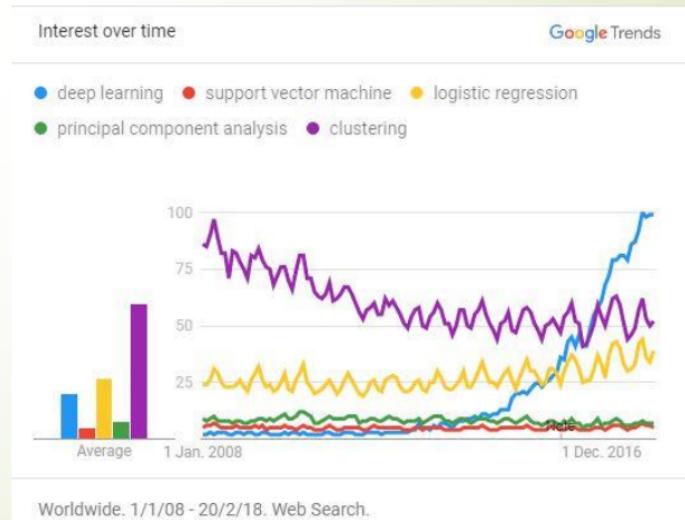
- A sub-field of machine learning for learning **representations** of data.
- Exceptionally effective at **learning patterns**.
- Deep learning algorithms attempt to learn (multiple levels of) representation by using a **hierarchy of multiple layers**
- If you provide the system **tons of information**, it begins to understand it and respond in useful ways.



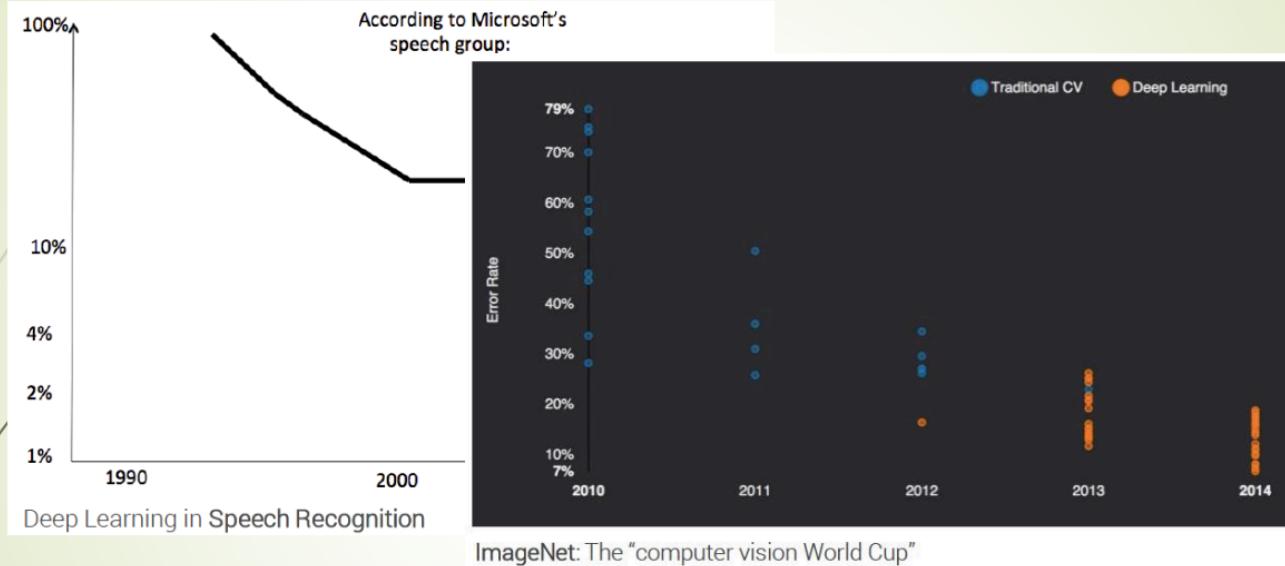
# Why is DL useful?

- Manually designed features are often **over-specified, incomplete** and take a **long time to design** and validate
- Learned Features are **easy to adapt, fast** to learn
- Deep learning provides a very **flexible**, (almost?) **universal**, learnable framework for representing world, visual and linguistic information.
- Can learn in both unsupervised and supervised ways
- Effective **end-to-end** joint system learning
- Utilize large amounts of training data

Around 2010, DL started to outperform other ML techniques, first in speech and vision, then in Natural Language Processing (NLP)



# State of the art in ...



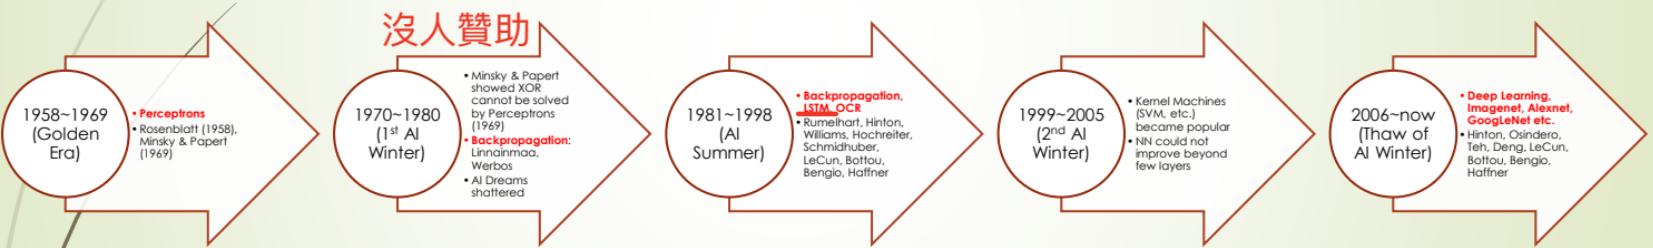
Several big improvements in recent years in NLP

- ✓ Machine Translation
- ✓ Sentiment Analysis
- ✓ Dialogue Agents
- ✓ Question Answering
- ✓ Text Classification ...

Leverage different levels of representation

- words & characters
- syntax & semantics

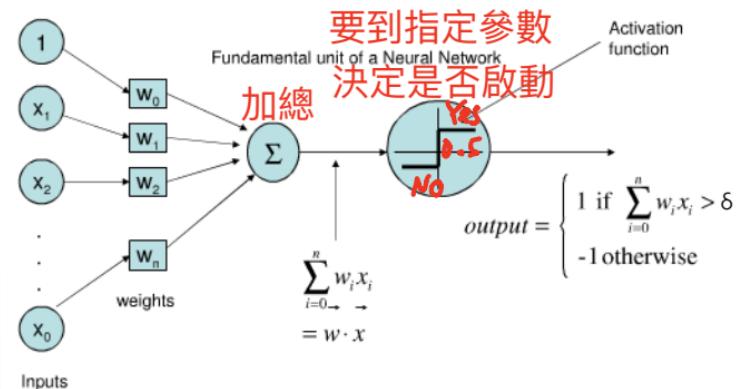
# History of neural networks & deep learning



# 1958~1969: Birth of AI and Golden Era

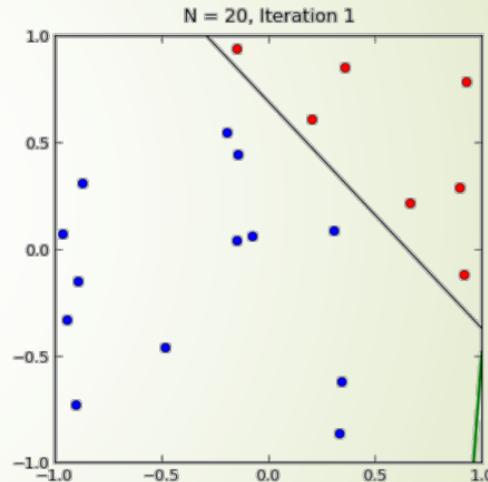
## -- Perceptrons 單層神經

- ▶ Rosenblatt proposed a machine for binary classifications
- ▶ Main idea
  - ▶ One weight  $w_i$  per input  $x_i$
  - ▶ Multiply weights with respective inputs and add bias  $w_0$
  - ▶ If result is larger than threshold  $\delta$ , return 1, otherwise 0.



# Training a perceptron

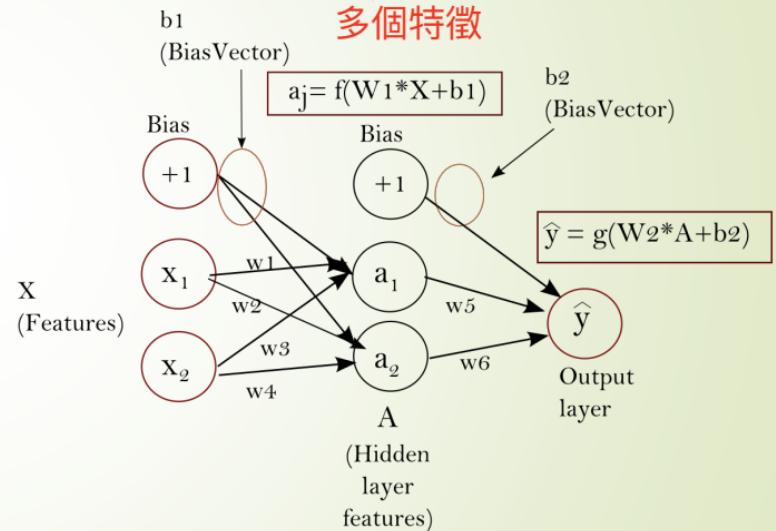
- ▶ **Learning Algorithm** by Rosenblatt
  - ▶ Initialize weights randomly
  - ▶ Take one sample  $x_i$  and predict  $y_i$
  - ▶ For erroneous predictions, update weights 錯誤要更改權重
    - ▶ If the output was  $\hat{y} = 0$  and  $y_i = 1$ , increase weights
    - ▶ If the output was  $\hat{y} = 1$  and  $y_i = 0$ , decrease weights
  - ▶ Repeat until no errors are made



Video (18'54''): <https://youtu.be/OVHc-7GYRo4>

# Multi-Layer Perceptron

- ▶ One perceptron = one decision
- ▶ **Question:** What about multiple decisions?
  - ▶ Eg. Digit classification
- ▶ **Answer:** Neural Network (NN) or Multi-Layer Perceptron (MLP)
  - ▶ Stack multiple perceptrons (neurons) into a single layer
  - ▶ Connect two or more layers by feeding output of one layer as input to the next layer



# 1970~1980: 1<sup>st</sup> AI Winter

- ▶ XOR cannot be solved by Perceptron (Minsky)
- ▶ Perceptron training method cannot be applied to Neural Networks
- ▶ Funding slushed, Neural Networks were damned
- ▶ **AI WINTER!!!**
- ▶ Dreams shattered!
- ▶ Some significant results
  - ▶ Backpropagation: training method for NN (1970, 1974)

# 1981~1998: AI Summer

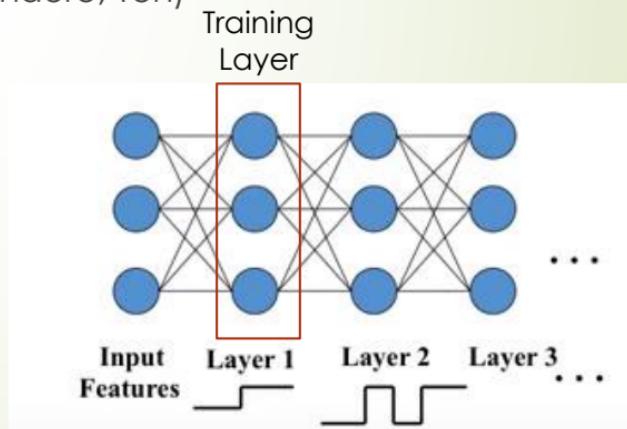
- ▶ XOR can be **solved** using Neural Networks (Multi-Layer Perceptron)
- ▶ **Backpropagation** (1986)
- ▶ Complicated NN architectures allowed
  - ▶ Temporal behavior, language, sequences could be learnt using “Recurrent Long-Short Term Memory (**LSTM**) Networks” in **1997** **連續性資料**
  - ▶ Digit recognition in cheques (OCR) solved using Convolutional Neural Network (**CNN**) in **1998** **物件偵測**

# 1999~2005: 2<sup>nd</sup> AI Winter

- ▶ Kernel Machines (e.g. **Support Vector Machines (SVM)**, etc.) became popular **分類方法**
  - ▶ Achieved similar accuracies
  - ▶ Included much fewer heuristics **探索的**
  - ▶ Nice proofs on generalization **一般化**
- ▶ Neural networks **could not improve** beyond a few layers
  - ▶ Lack of **processing power** (No GPUs)
  - ▶ Lack of **data** (No big, annotated datasets)
  - ▶ Overfitting (Models could not generalize) **過度訓練**
  - ▶ Vanishing gradients ( $0.1 * 0.1 * 0.1 * \dots * 0.1 = 0.000000000001$ , too small for learning)
- ▶ AI community turned away from Neural Networks

# 2006~now: Thaw of AI Winter

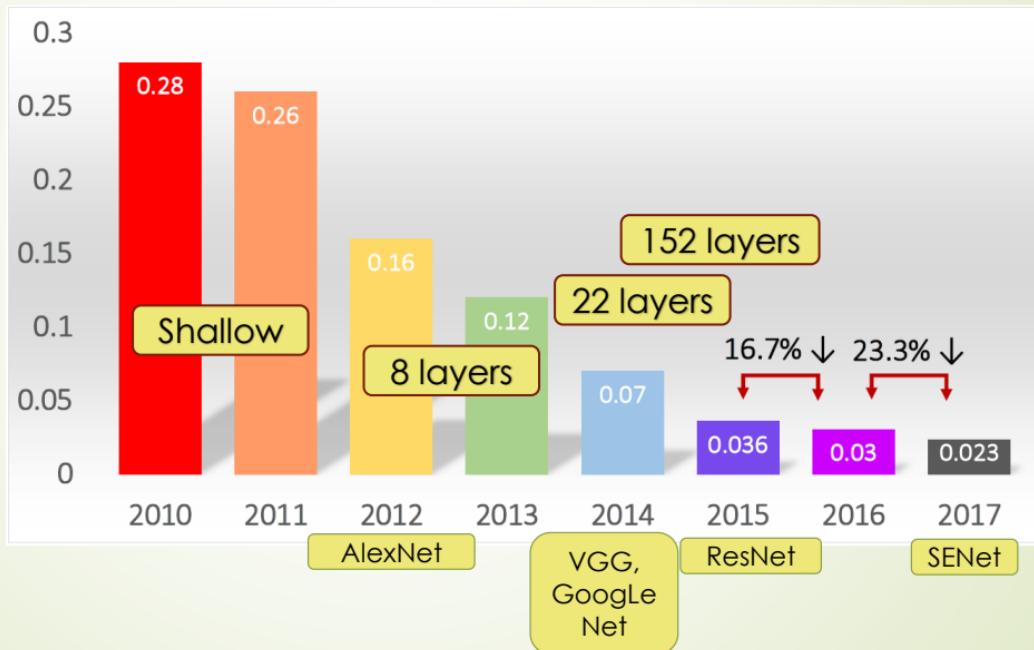
- ▶ Questions: 訓練層
  - ▶ Are 1-2 hidden layers the best NN can do?
  - ▶ Or, is it the learning algorithm not really mature?
- ▶ Deep Learning (2006, Hinton, Osindero, Teh)
- ▶ Layer-by-layer training
  - ▶ Per-layer trained parameters initialize further training using contrastive divergence  
對比的 分離



# Deep Learning is here ...

- ▶ **ImageNet** dataset (Deng et al, 2009)
  - ▶ Collected images for each term of Wordnet (100,000 classes)
  - ▶ Tree of concepts organized hierarchically
    - ▶ "Ambulance", "Dalmatian dog", "Egyptian cat", ...
- ▶ **Imagenet Large Scale Visual Recognition Challenge** (ILSVRC)
  - ▶ 1 million images
  - ▶ 1,000 classes
  - ▶ Top-5 and top-1 error measured
    - ▶ Errors reduced drastically in the past 8 years (2010~2017): 28.2% → 2.3%

# ImageNet Large Scale Visual Recognition Challenge (ILSVRC)



# Some FUN now ...



- ▶ The Neural Network ZOO
  - ▶ Graphical notations for all kinds of neural networks
  - ▶ <http://www.asimovinstitute.org/neural-network-zoo/>
- ▶ A Neural Network Playground
  - ▶ An online interactive way to play with different network architectures
  - ▶ <http://playground.tensorflow.org>
- ▶ 8 Inspirational Applications of Deep Learning
  - ▶ Very interesting applications of deep learning
  - ▶ <http://machinelearningmastery.com/inspirational-applications-deep-learning/>

# (1/8) Automatic Colorization of B&W Images

- ▶ Large Convolutional Neural Networks (CNN) 物件辨識
- ▶ Website
  - ▶ <http://richzhang.github.io/colorization/>
- ▶ Video (5 s)
  - ▶ <http://whattogive.com/videoColourization/>



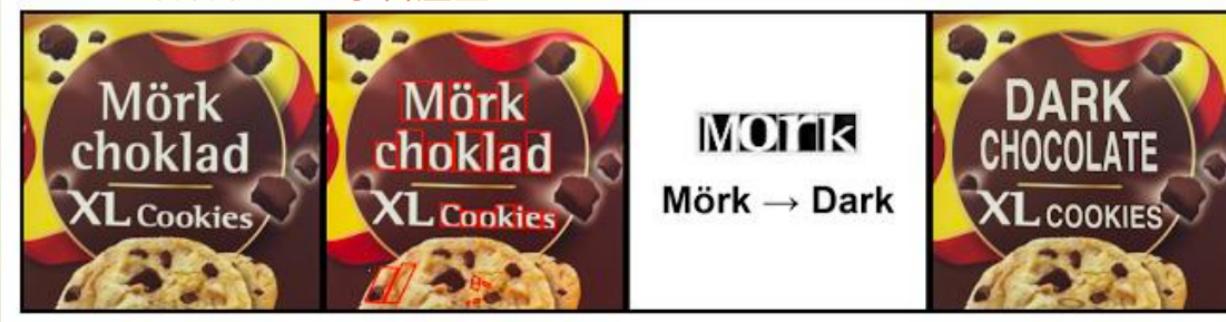
## (2/8) Automatically Adding Sounds

- ▶ Two types of NN 時序性資料
  - ▶ Large CNN for images
  - ▶ Large Long Short-Term Memory (LSTM) Recurrent Neural Networks (RNN) for sound
- ▶ News
  - ▶ <http://news.mit.edu/2016/artificial-intelligence-produces-realistic-sounds-0613>
- ▶ Visually Indicated Sounds (MIT)
  - ▶ <http://vis.csail.mit.edu/>
- ▶ Video (2.54 s)
  - ▶ <https://youtu.be/0FW99AQmMc8>

# (3/8) Automatic Machine Translation

- ▶ Automatic Translation of Text
  - ▶ Large Long Short-Term Memory (LSTM) Recurrent Neural Networks
- ▶ Automatic Translation of Images
  - ▶ CNN + LSTM RNN

物件偵測 學習產生

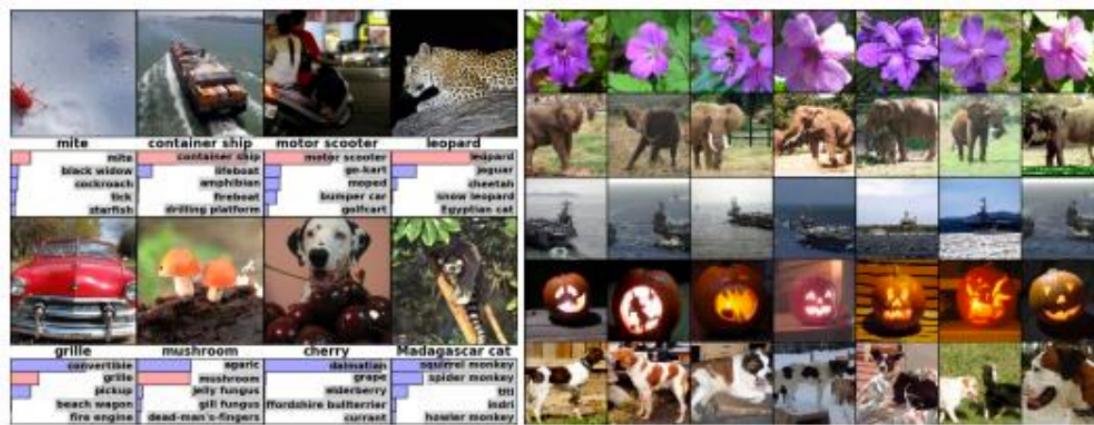


# How Google Translate squeezes deep learning onto a phone!

- ▶ Demonstration

# (4/8) Object Classification & Detection in Photographs

- ▶ Large deep CNN
- ▶ Paper on ImageNet Classification
  - ▶ <http://www.cs.toronto.edu/~fritz/absps/imagenet.pdf>



# ConvNetJS: CIFAR-10 Demo

- ▶ ConvNetJS: CIFAR-10 Demo
  - ▶ <http://cs.stanford.edu/people/karpathy/convnetjs/demo/cifar10.html>
- ▶ Clarifai: 10,000 images
  - ▶ <https://www.clarifai.com/>

## (5/8) Automatic Handwriting Generation

- ▶ Usage: used with forensic analysis

Machine learning Mastery

Madine Learning Mastery

Mtachne Learning Mastery

# Automatic Handwriting Generation DEMO

► <http://www.cs.toronto.edu/~graves/handwriting.html>

# (6/8) Automatic Text Generation

- ▶ Large RNN
- ▶ Code on Github
  - ▶ <https://github.com/karpathy/char-rnn>
- ▶ Paul Graham generator
- ▶ Shakespeare
- ▶ Wikipedia
- ▶ Algebraic Geometry (LaTeX)
- ▶ Linux source code
- ▶ Generating Baby Names

PANDARUS:

Alas, I think he shall be come approached and the day  
When little strain would be attain'd into being never fed,  
And who is but a chain and subjects of his death,  
I should not sleep.

Second Senator:

They are away this miseries, produced upon my soul,  
Breaking and strongly should be buried, when I perish  
The earth and thoughts of many states.

DUKE VINCENTIO:

Well, your wit is in the care of side and that.

Second Lord:

They would be ruled after this chamber, and  
my fair nues begun out of the fact, to be conveyed,  
Whose noble souls I'll have the heart of the wars.

Clown:

Come, sir, I will make did behold your worship.

VIOLA:

I'll drink it.

# Automatic Music Synthesis

- ▶ Using large RNN
- ▶ Authors comments:
  - ▶ This track was made using a RNN. Fed 500 mb guitar tabs in ASCII. It writes the tabs out in ASCII, I imported into GuitarPro, recorded the output, imported that into FL Studio, added some filters and a drum loop and got this. The notes and rhythms themselves are totally unedited.
- ▶ Music (5:20):
  - ▶ <https://soundcloud.com/optometrist-prime/recurrence-music-written-by-a-recurrent-neural-network>

## (7/8) Automatic Image Caption Generation

- ▶ Large CNN → Object Detection
- ▶ Large LSTM RNN → Caption Text Generation
- ▶ Deep Visual-Semantic Alignments for Generating Image Descriptions
  - ▶ <http://cs.stanford.edu/people/karpathy/deepimagesent/>

# Automatic Image Captioning

## ► Demo

- <http://cs.stanford.edu/people/karpathy/doingimagecaptioningandmore/>
- 

"man in black shirt is playing guitar."



"construction worker in orange safety vest is working on road."



"two young girls are playing with lego toy."
- 

"girl in pink dress is jumping in air."



"black and white dog jumps over bar."



"young girl in pink shirt is swinging on swing."

## (8/8) Automatic Game Playing

- ▶ Vision
- ▶ Decision Making
- ▶ Etc.

# Game Playing Demo

► <https://youtu.be/TmPfTpjtdgg>

# AlphaGo

- ▶ AlphaGo by Google Deepmind Team
  - ▶ <https://deepmind.com/research/alphago>
- ▶ News:
  - ▶ AlphaGo program played against **Lee Sedol (18 times world champion)** in a tournament that took a week to play, five matches between March 9-15, 2016 at the Four Seasons Hotel in Seoul, South Korea.
  - ▶ It became the first program to beat a professional Go player, and in doing so showed an advancement of artificial intelligence that took the world by surprise.
  - ▶ Press Conference
    - ▶ <https://youtu.be/yCALyQRN3hw?t=20728>

# AlphaGo in the news!



# AlphaGo



4:1

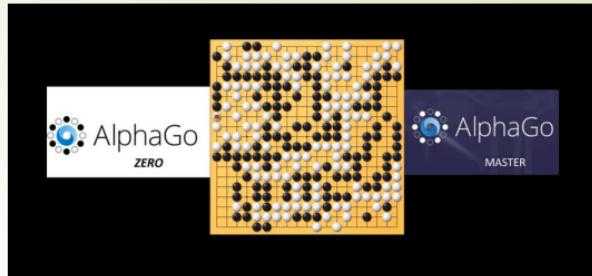
5:0

# Alpha Go

- ▶ Used **general purpose** algorithms
  - ▶ Not a set of handcraft rules
  - ▶ Very different from Deep Blue (expert system)
- ▶ Modular system combining **planning and pattern recognition**
- ▶ Think like **human**
- ▶ **Reinforcement** and deep learning

# AlphaGo Zero

- ▶ Alpha Go
  - ▶ Learn from large amounts of HUMAN PLAYS
  - ▶ Defeated the World Champion Lee Sedol from Korea
- ▶ AlphaGo Zero
  - ▶ **Learning by itself (No human data)!**
  - ▶ Defeated all previous versions of Alpha Go
  - ▶ Used **less computation** than previous versions of Alpha Go
  - ▶ Defeated the version of Alpha Go that defeated the World Champion (**result: 100:1**)
- ▶ About: <https://www.youtube.com/watch?v=tXIM99xPQC8>
- ▶ Nature Paper: <https://www.nature.com/articles/nature24270>
- ▶ AlphaGo Zero vs. AlphaGo Master Game1: <https://youtu.be/-Wh4CfsWDyM>
- ▶ AlphaGo Zero vs. AlphaGo Master Game2: <https://youtu.be/xOVwmCOX7S4>



# Alpha Zero

- ▶ Not only Go!
  - ▶ Also **Chess** and **Shogi**
- ▶ Self-play using **5,000 first-generation TPUs** to generate games, **64 second-generation TPUs** to train the NN, all in parallel.
- ▶ Differences between AlphaGo Zero
  - ▶ Alpha Zero has **hard-coded rules** for setting search hyperparameters
  - ▶ NN is **updated continually**
  - ▶ Did not take advantage of **symmetry** (AlphaGo Zero used symmetry in Go)
  - ▶ Chess can end in a **draw**, AlphaZero takes draws into account.

# Techniques required for DL

- ▶ Python Programming Language
- ▶ Deep Learning Frameworks
- ▶ Calculus

# Python programming language

- ▶ Why Python?
- ▶ More Information on Python
- ▶ Installing Python
- ▶ Jupyter Notebooks
- ▶ Examples of running Python
- ▶ Running Python in Jupyter Notebooks

# Python programming language



Guido van  
Rossum

- ▶ Conceived in late 1980's
- ▶ Implementation started in December 1989
- ▶ Creator: Guido van Rossum at CWI in the Netherlands (2005-2012 with Google, now with Dropbox)
- ▶ As successor to ABC programming language, with exception handling and for Amoeba OS
- ▶ Python named after BBC TV show Monty Python's Flying Circus



# Why Python?

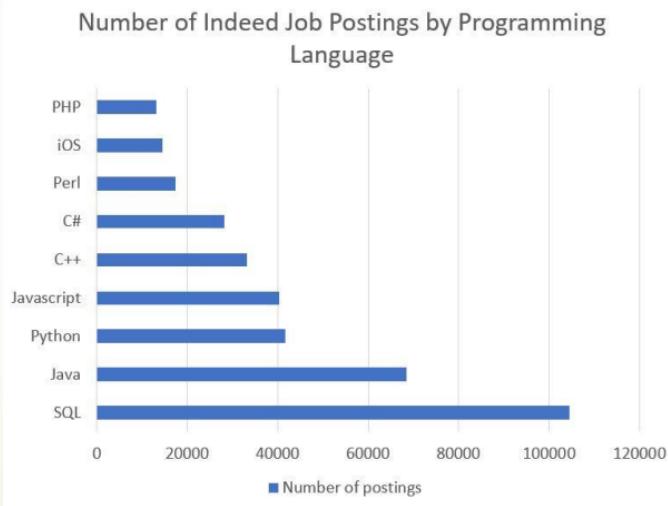
## ► 2017 Top Programming Languages by IEEE Spectrum



Source: <https://spectrum.ieee.org/computing/software/the-2017-top-programming-languages>

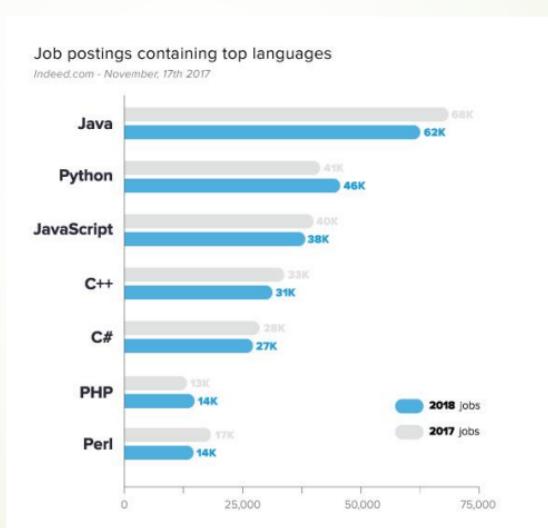
# Why Python?

► **2017 Top Programming Languages by Number of Job Postings** by Coding Dojo Blog



# Why Python?

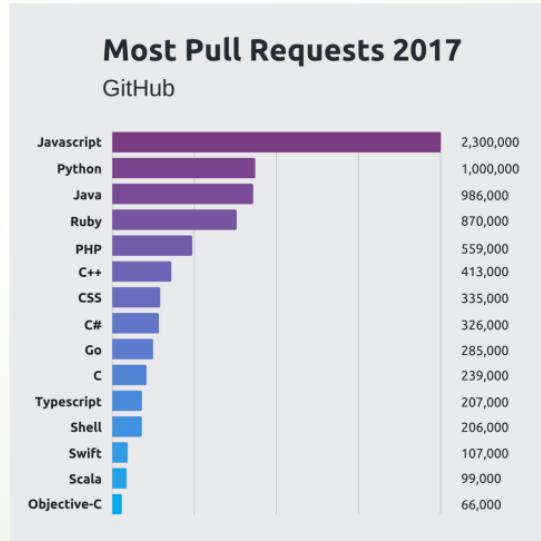
- The **7 Most-in-Demand Programming Languages of 2018** by Coding Dojo Blog



Source: <http://www.codingdojo.com/blog/7-most-in-demand-programming-languages-of-2018/>

# Why Python?

## ► Most Pull Requests on GitHub in 2017



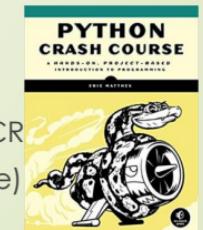
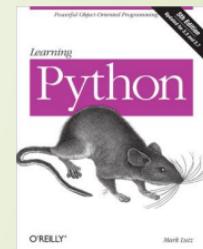
Source: <https://stackify.com/popular-programming-languages-2018/>

# Why Python?

- ▶ **It is Popular:** 1 million Python users
- ▶ **It has NICE features:** Object-oriented, open-source, portable
- ▶ **It is POWERFUL:** dynamic typing, built-in types and tools, libraries, 3<sup>rd</sup> party utilities (Numeric, Numpy, SciPy, etc.), automatic memory management
- ▶ **It has W I D E applications**
  - ▶ System/GUI programming, Internet Scripting, Component Integration, Database, Gaming, Images, XML, Robot, etc.
  - ▶ Google: uses Python in web search system (Python's creator employed)
  - ▶ Intel, Cisco, HP, Seagate, Qualcomm, IBM: use Python for hardware testing
  - ▶ YouTube: video sharing services written in Python mostly

# More Information on Python

- ▶ <http://python.org/>
  - ▶ Documentation, tutorials, beginners guide, core distribution, etc.
- ▶ Books
  - ▶ <https://wiki.python.org/moin/PythonBooks>
  - ▶ <http://learning-python.com/about-lp5e.html>
    - ▶ Learning Python by Mark Lutz, O'Reilly
- ▶ Online Python Popular Code Recipes
  - ▶ <http://code.activestate.com/recipes/langs/python/> (author site)
  - ▶ <http://www.dsf.unica.it/~fiore/LearningPython.pdf> (e-book)
- ▶ Python Crash Course (e-book)
  - ▶ <http://ap-n.us/books/Programming/Python%20Crash%20Course.pdf>
- ▶ Videos
  - ▶ <https://code.tutsplus.com/courses/introduction-to-python> (Learn from SCR)
  - ▶ [https://www.tutorialspoint.com/python\\_online\\_training/index.asp](https://www.tutorialspoint.com/python_online_training/index.asp) (Course)



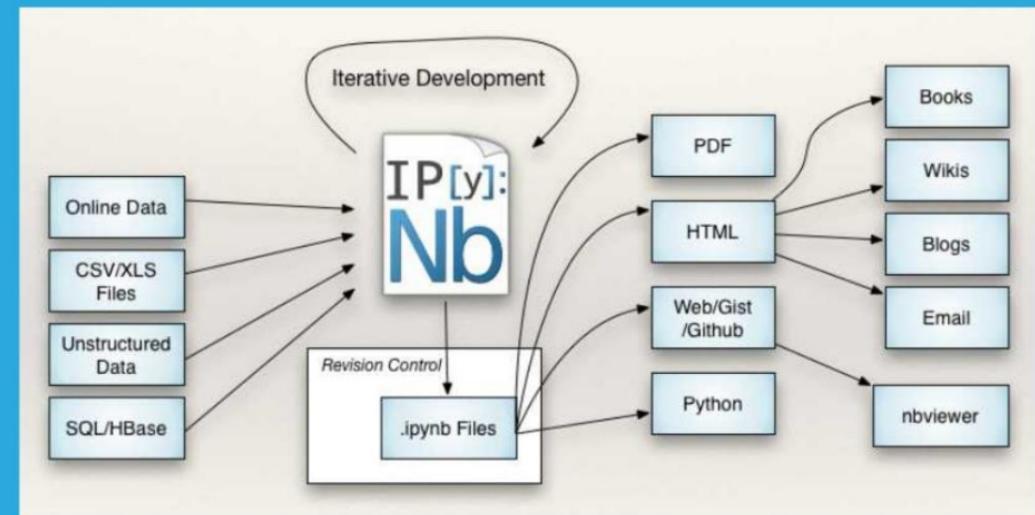
# Installing Python

- ▶ <https://www.python.org/downloads/>
- ▶ Mainly TWO versions:
  - ▶ Python2 (currently 2.7.14)
  - ▶ Python3 (currently 3.6.4)
- ▶ Platforms supported: Windows, Linux/UNIX, Mac OS X, Others
- ▶ Beginners' guide
  - ▶ <https://wiki.python.org/moin/BeginnersGuide/Programmers>
- ▶ Which IDE to use for Python?
  - ▶ <https://docs.google.com/spreadsheets/d/1I3x94P55qoxqYbq5GosWQ7lonZ4vR-4ZyCalmiVmCSk/pubhtml> (there are lots of IDE, among which IDLE has a GUI)
  - ▶ In this course, you can use any as long as you can submit your work assignments
  - ▶ For labs, we will use Jupyter Notebooks

# Jupyter Notebooks

- ▶ An open-source web application for **interactive** data science and scientific computing across **40 programming languages including Python, R, Julia and Scala ...**
- ▶ Combines LIVE code and NARRATIVE text, equations, images, video, and visualizations.
- ▶ Can be shared on GitHub, Dropbox, and Jupyter Notebook Viewer
- ▶ Leverage big-data tools such as Apache Spark, from Python, R, and Scala. Explore data with pandas, scikit-learn, ggplot2, dplyr, etc.

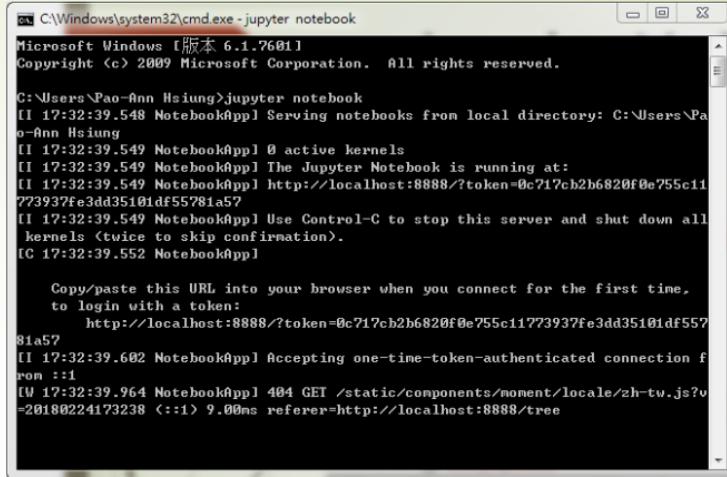
# Jupyter Notebooks



Credit: Joshua Barrat (via F. Perez)

# Jupyter Notebooks

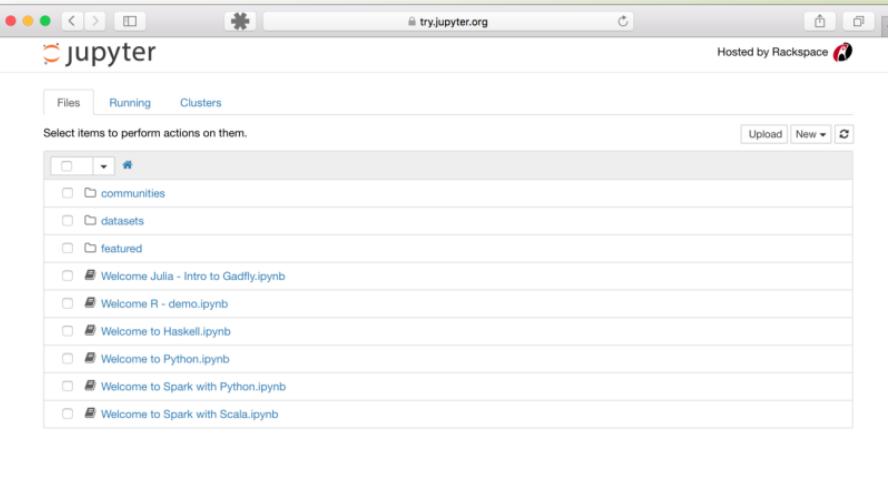
- ▶ Main page: <http://jupyter.org>
- ▶ Try it in your browser:  
<https://try.jupyter.org/>
- ▶ Install Jupyter:  
<http://jupyter.org/install.html>
- ▶ Running Jupyter Notebook:  
<https://jupyter.readthedocs.io/en/latest/running.html#running>



```
C:\Windows\system32\cmd.exe - jupyter notebook
Microsoft Windows [版本 6.1.7601]
Copyright © 2009 Microsoft Corporation. All rights reserved.

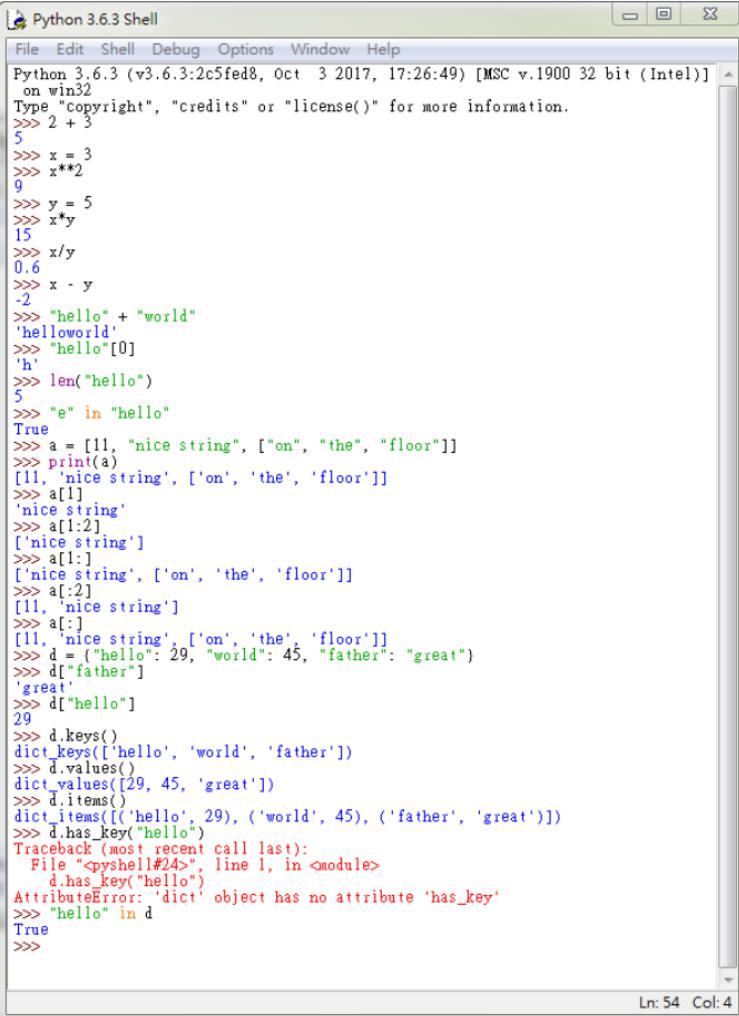
C:\Users\Pao-An Hsiung>jupyter notebook
[I 17:32:39.548 NotebookApp] Serving notebooks from local directory: C:\Users\Pao-An Hsiung
[I 17:32:39.549 NotebookApp] 0 active kernels
[I 17:32:39.549 NotebookApp] The Jupyter Notebook is running at:
[I 17:32:39.549 NotebookApp] http://localhost:8888/?token=0c717cb2b6820f0e755c11723937fe3dd35101df55781a57
[I 17:32:39.549 NotebookApp] Use Control-C to stop this server and shut down all
kernels <twice to skip confirmation>.
[IC 17:32:39.552 NotebookApp]

Copy/paste this URL into your browser when you connect for the first time,
to login with a token:
    http://localhost:8888/?token=0c717cb2b6820f0e755c11723937fe3dd35101df55781a57
[I 17:32:39.602 NotebookApp] Accepting one-time-token-authenticated connection from ::1
[EW 17:32:39.964 NotebookApp] 404 GET /static/components/moment/locale/zh-tw.js?v=20180224173238 ::1 9.00ms referer=http://localhost:8888/tree
```



# Examples of running Python

- ▶ Data types
  - ▶ Numbers, Strings, Lists, Dictionaries
- ▶ List items and slicing
- ▶ Dictionary keys and values



The screenshot shows the Python 3.6.3 Shell interface. The window title is "Python 3.6.3 Shell". The menu bar includes File, Edit, Shell, Debug, Options, Window, and Help. The main area displays Python code and its output. The code demonstrates various operations on lists and dictionaries, such as concatenation, indexing, slicing, and key-value access.

```
Python 3.6.3 (v3.6.3:2c5fed8, Oct  3 2017, 17:26:49) [MSC v.1900 32 bit (Intel)]
Type "copyright", "credits" or "license()" for more information.
>>> 2 + 3
5
>>> x = 3
>>> x**2
9
>>> y = 5
>>> x*y
15
>>> x/y
0.6
>>> x - y
-2
>>> "hello" + "world"
'helloworld'
>>> "hello"[0]
'h'
>>> len("hello")
5
>>> "e" in "hello"
True
>>> a = [11, "nice string", ["on", "the", "floor"]]
>>> print(a)
[11, 'nice string', ['on', 'the', 'floor']]
>>> a[1]
'nice string'
>>> a[1:2]
['nice string']
>>> a[1:]
['nice string', ['on', 'the', 'floor']]
>>> a[-2]
[11, 'nice string']
>>> a[:]
[11, 'nice string', ['on', 'the', 'floor']]
>>> d = {"hello": 29, "world": 45, "father": "great"}
>>> d["father"]
'great'
>>> d["hello"]
29
>>> d.keys()
dict_keys(['hello', 'world', 'father'])
>>> d.values()
dict_values([29, 45, 'great'])
>>> d.items()
dict_items([('hello', 29), ('world', 45), ('father', 'great')])
>>> d.has_key('hello')
Traceback (most recent call last):
  File "<pyshell#24>", line 1, in <module>
    d.has_key("hello")
AttributeError: 'dict' object has no attribute 'has_key'
>>> "hello" in d
True
>>>
```

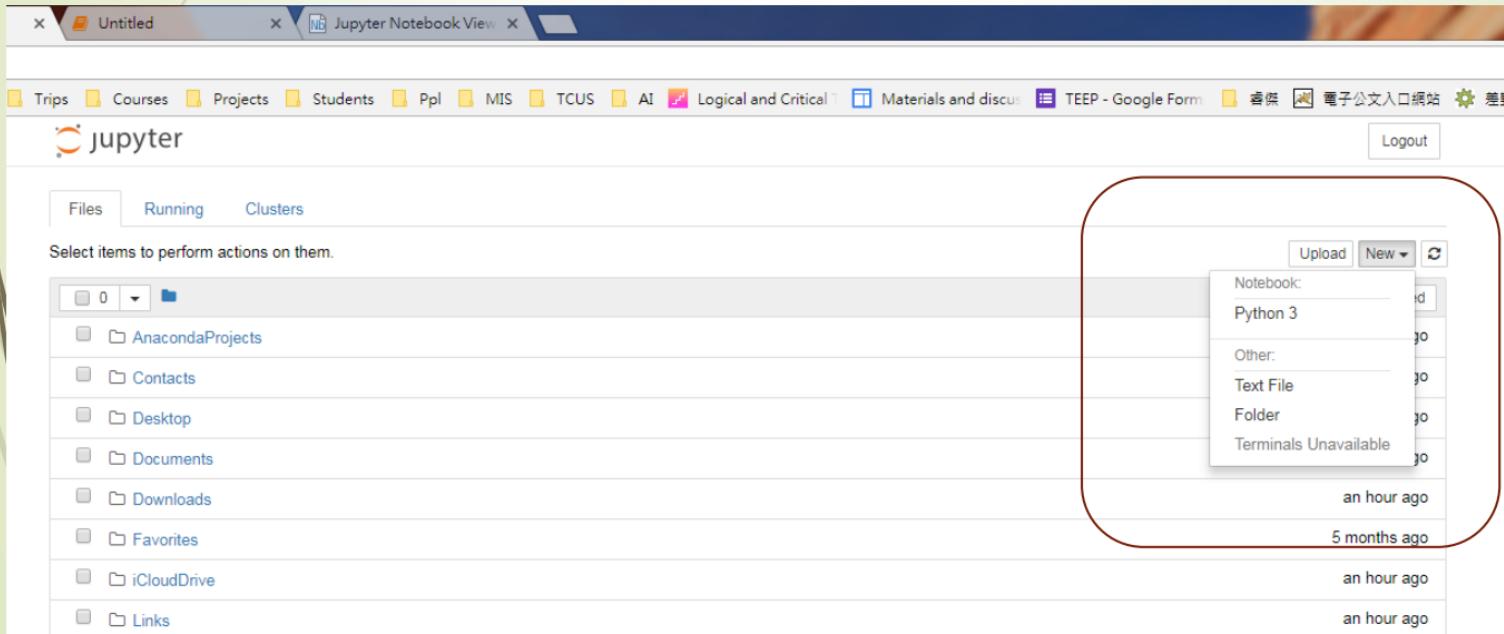
# Running Python in Jupyter Notebooks

- ▶ A notebook is made up of a number of **cells**
  - ▶ Each cell can contain **Python code, text, images, etc.**
- ▶ You can **execute** a cell by clicking on it and pressing **Shift-Enter**.
  - ▶ The **output** of the cell will be displayed **below**.
- ▶ By convention, cells are to be run from top to bottom.
  - ▶ Failing to execute some cells or executing cells out of order can result in errors.
- ▶ More questions? Check out the following Jupyter Notebook Tutorial:
  - ▶ <https://www.datacamp.com/community/tutorials/tutorial-jupyter-notebook>

# Examples of running Python in Notebook

- ▶ Several Examples

- ▶ <https://github.com/jakevdp/PythonDataScienceHandbook/tree/master/notebooks>



nb?kernel\_name=python3

Trips Courses Projects Students Ppl MIS TCUS AI Logical and Critical Materials and discuss TEEP - Google Form 電子公文入口網站 Logout

jupyter Untitled Last Checkpoint: 4 minutes ago (unsaved changes)

File Edit View Insert Cell Kernel Help Trusted Python 3

In [1]: `2 + 3`

Out[1]: 5

In [2]: `x = 2`

In [3]: `x**2`

Out[3]: 4

In [4]: `y = 5`

In [5]: `x*y`

Out[5]: 10

In [6]: `x/y`

Out[6]: 0.4

In [7]: `x - y`

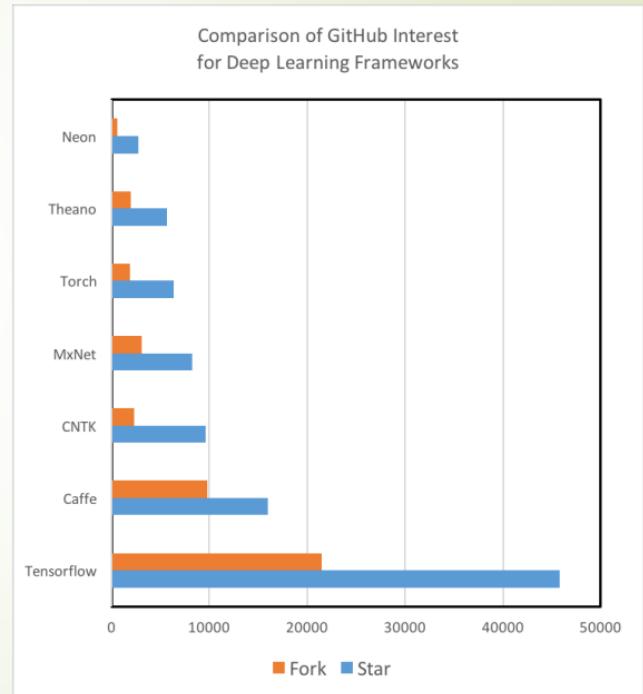
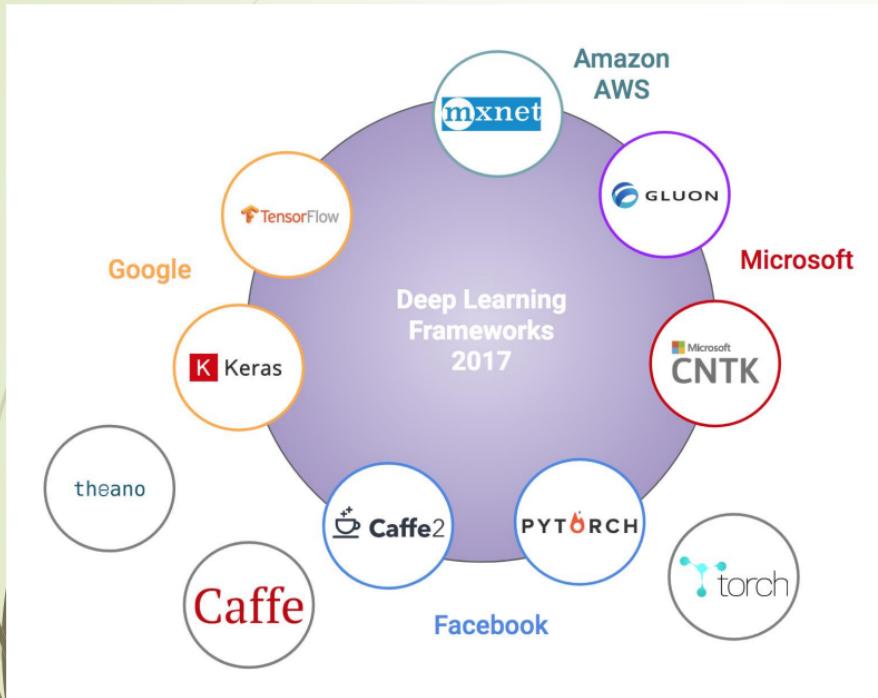
Out[7]: -3

In [8]: `a = ["hello", "world", "father"]`

In [9]: `a[0] + " " + a[1] + "!"`

Out[9]: 'hello world!'

# Deep learning frameworks

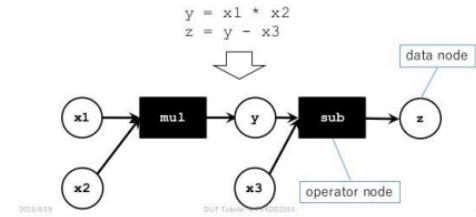


# Deep learning frameworks

- ▶ Easily build big computational graphs
- ▶ Easily compute gradients in computational graphs
- ▶ Run it all efficiently on GPU

## Neural Network as a Computational Graph

- In simplest form, NN is represented as a computational graph (CG) that is a stack of bipartite DAGs (Directed Acyclic Graph) consisting of **data nodes** and **operator nodes**.



[https://en.wikipedia.org/wiki/Comparison\\_of\\_deep\\_learning\\_software](https://en.wikipedia.org/wiki/Comparison_of_deep_learning_software)

[http://cs231n.stanford.edu/slides/2017/cs231n\\_2017\\_lecture8.pdf](http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture8.pdf)

<https://towardsdatascience.com/a-survey-of-deep-learning-frameworks-43b88b11af34>

# Comparisons in modeling capabilities

<https://www.kdnuggets.com/2017/03/getting-started-deep-learning.html>

	Languages	Tutorials and training materials	CNN modeling capability	RNN modeling capability	Architecture: easy-to-use and modular front end	Speed	Multiple GPU support	Keras compatible
Theano	Python, C++	++	++	++	+	++	+	+
Tensor-Flow	Python	+++	+++	++	+++	++	++	+
Torch	Lua, Python (new)	+	+++	++	++	+++	++	
Caffe	C++	+	++		+	+	+	
MXNet	R, Python, Julia, Scala	++	++	+	++	++	+++	
Neon	Python	+	++	+	+	++	+	
CNTK	C++	+	+	+++	+	++	+	

# Comparisons in code

## Numpy

```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```

## TensorFlow

```
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3, 4

with tf.device('/gpu:0'):
    x = tf.placeholder(tf.float32)
    y = tf.placeholder(tf.float32)
    z = tf.placeholder(tf.float32)

    a = x * y
    b = a + z
    c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
    out = sess.run([c, grad_x, grad_y, grad_z],
                  feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```

## PyTorch

```
import torch
from torch.autograd import Variable

N, D = 3, 4

x = Variable(torch.randn(N, D).cuda(),
             requires_grad=True)
y = Variable(torch.randn(N, D).cuda(),
             requires_grad=True)
z = Variable(torch.randn(N, D).cuda(),
             requires_grad=True)

a = x * y
b = a + z
c = torch.sum(b)

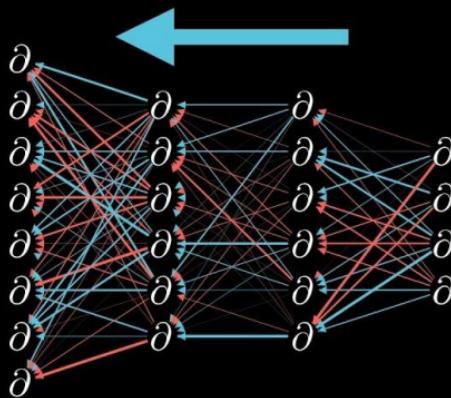
c.backward()

print(x.grad.data)
print(y.grad.data)
print(z.grad.data)
```

# Calculus

- ▶ Partial derivatives
- ▶ Chain rule

## Backpropagation calculus



<https://www.youtube.com/watch?v=tleHLnjs5U8>

# What we learnt?

- ▶ Bonus points for correct and accurate answers to the following questions:
  - ▶ What is **machine** learning?
  - ▶ What is the difference between **machine** and **deep** learning?
  - ▶ Can CNN and RNN be **combined** into one neural network architecture? If yes, for what kind of **applications**?
  - ▶ Given the following Python code:

```
>>> def f4(x,a=1):  
...     return a*x**2  
...  
>>>
```

- ▶ What is the value of **f4(2)**?
- ▶ What is the value of **f4(2, 2)**?

# What to do now?

- ▶ Install Python
- ▶ Install Jupyter
- ▶ Play with Jupyter Notebooks and Python
- ▶ Next class on March 6 (Tuesday) 13:15 will be in the **COMPUTER ROOM R341 in the INNOVATION BUILDING!!!**
  - ▶ We will have a lab in the computer room for only 75 minutes
  - ▶ The second class 14:45~16:00 will still be conducted in room EA204 here!