

# Introduction to Deep Learning

## Chapter 7: Deep Generative Models

Pao-Ann Hsiung

National Chung Cheng University

# Contents

- Introduction
- Autoencoder
- Generative Adversarial Networks (GAN)
- Variants of GAN
- Applications of GAN
- Implementation of GAN
- Issues and Conclusions

# Contents

- Introduction
- Autoencoder
- Generative Adversarial Networks (GAN)
- Variants of GAN
- Applications of GAN
- Implementation of GAN
- Issues and Conclusions

# Introduction

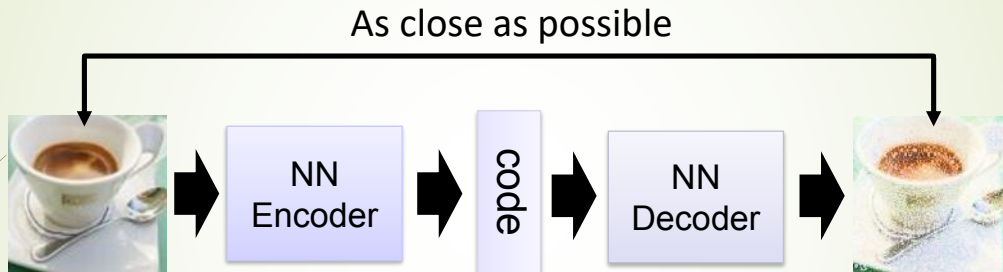
## ► Why **generation**?

- Classification (class output) or regression (scalar output) is limited in scope of applications
- Structured outputs are often required for structured inputs
  - Examples: language translation, text to image, autonomous driving images, etc.


# Contents

- Introduction
- Autoencoder
- Generative Adversarial Networks (GAN)
- Variants of GAN
- Applications of GAN
- Implementation of GAN
- Issues and Conclusions

# Autoencoder




Randomly  
generate a vector  
as code



# Autoencoder with 3 fully connected layers

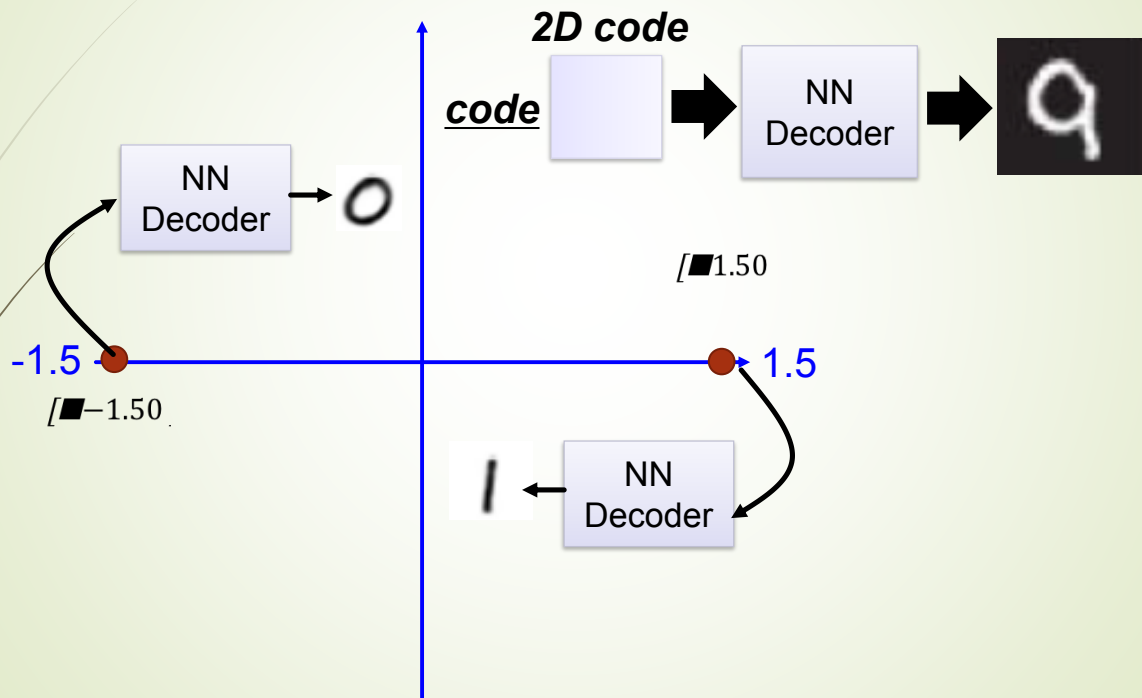
Training: `model.fit(X,X)`

Cost function:  $\sum_{k=1..N} (x_k - x'_k)^2$




Large → small, learn to compress

# Auto-encoder






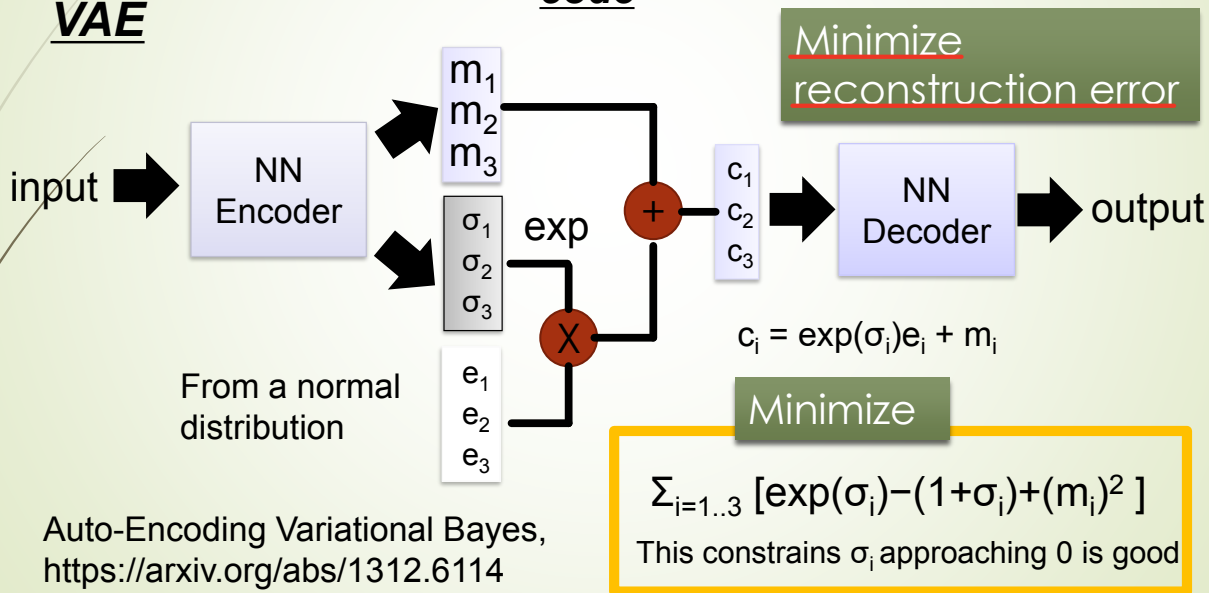
# Auto-encoder



# Auto-encoder

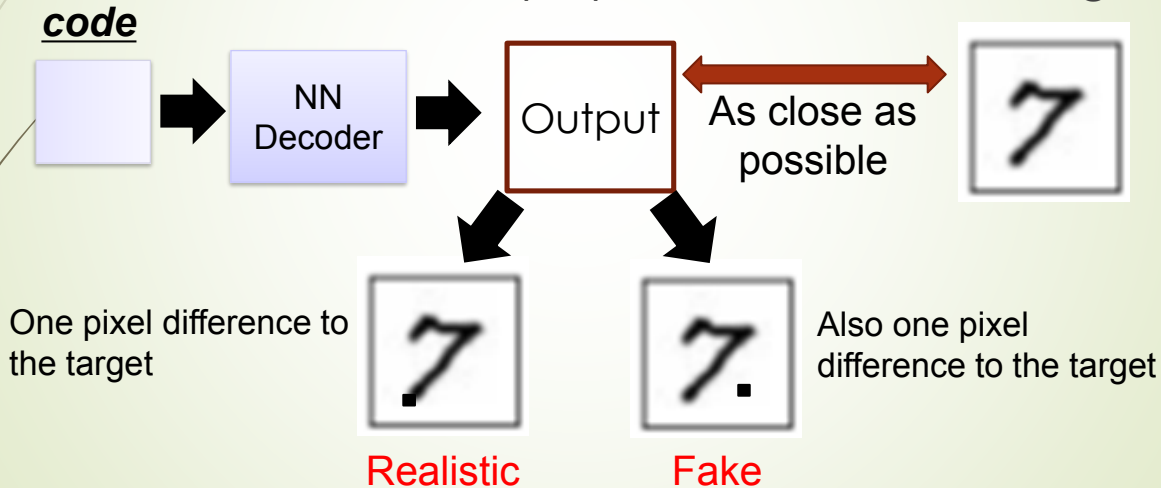


## VAE



# Problems of VAE

➡ It does not really try to simulate real images



VAE treats these the same

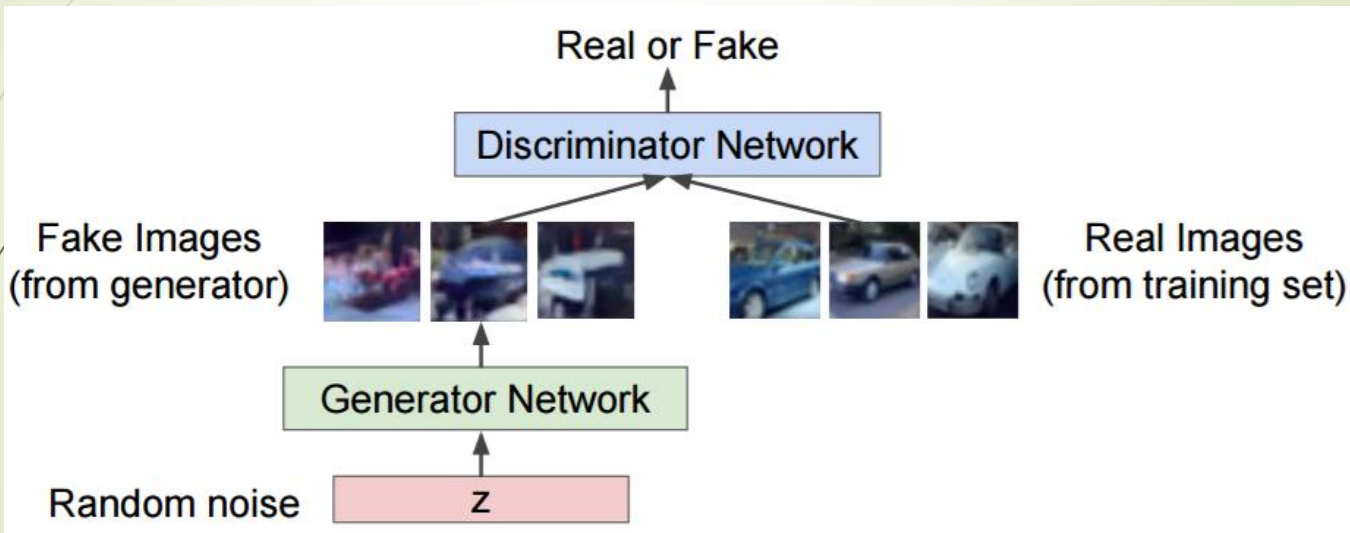
# Contents

- Introduction
- Autoencoder
- **Generative Adversarial Networks (GAN)**
- Variants of GAN
- Applications of GAN
- Implementation of GAN
- Issues and Conclusions

## What are GANs?

- System of two neural networks competing against each other in a zero-sum game framework.
  - Generator and Discriminator
    - dueling each other
- First introduced by Ian Goodfellow *et al.* in 2014.

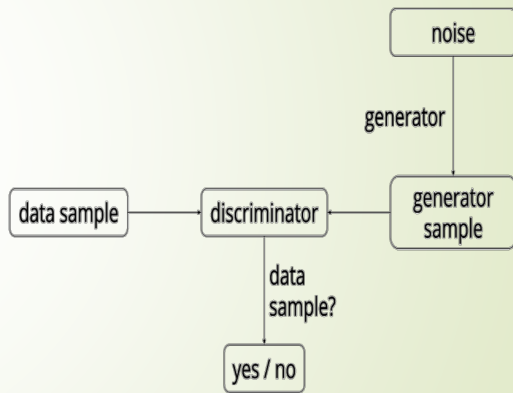
# GAN: Generator & Discriminator



# GAN: Generator & Discriminator

Given a set of target samples, the **Generator** tries to produce samples that can **fool the Discriminator** into believing they are real. The **Discriminator** tries to **distinguish real (target) samples** from **fake (generated) samples**.

After iterative training, we eventually end up with a **Generator that is really good at generating samples similar to the target samples**.



Overview of GANs

Source: <https://ishmaelbelghazi.github.io/ALI>

## Discriminative Models

- ▶ A **discriminative** model learns a function that maps the **input data (x)** to some desired **output class label (y)**.
- ▶ In probabilistic terms, they directly learn the conditional distribution  $P(y | x)$ .



# Generative Models


- ▶ A **generative** model tries to learn the joint probability of the input data and labels simultaneously i.e.  $P(x,y)$ .
- ▶ Potential to understand and explain the underlying structure of the input data even when there are no labels.

## How GANs are being used?

- Applied for modelling natural images.
- Performance is fairly good in comparison to other generative models.
- Useful for unsupervised learning tasks.

## Why GANs?

- ▶ Use a latent code.
- ▶ Asymptotically consistent (unlike variational methods) .
- ▶ No Markov chains needed.
- ▶ Often regarded as producing the best samples.



## How to train GANs?

**Generative與discriminative都需要train**

- Objective of generative network - increase the error rate of the discriminative network.
- Objective of discriminative network - decrease binary classification loss.
- Discriminator training - backprop from a binary classification loss.
- Generator training - backprop the **否定** negation of the binary classification loss of the discriminator.

# Loss Functions

$$\mathcal{L}(\hat{x}) = \min_{x \in data} (x - \hat{x})^2$$

**Generator**

$$D_G^*(x) = \frac{p_{data}(x) \text{ 大約 } 1}{p_{data}(x) + p_g(x) \text{ 大約 } 0}$$

**Discriminator**



Generated bedrooms. Source: "Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks" <https://arxiv.org/abs/1511.06434v2>



# Contents


- Introduction
- Autoencoder
- Generative Adversarial Networks (GAN)
- Variants of GAN
- Applications of GAN
- Implementation of GAN
- Issues and Conclusions



# Variants of GAN

- Look up this website for all current variants of GAN:
- <https://github.com/hindupuravinash/the-gan-zoo>
- To name a few ...
  - 3D-ED-GAN, 3D-GAN, 3d-IWGAN, 3D-PhyNet, 3D-RecGAN, ABC-GAN, AC-GAN, acGAN, ACGAN, AdaGAN, Adaptive GAN, AdvEntuRe, AdvGAN, AE-GAN, AE-OT, AEGAN, AF-DCGAN, AffGAN, AIM, AL-CGAN, ALI, AlignGAN, AlphaGAN, AM-GAN, AmbientGAN, AMC-GAN, AnoGAN, APD, APE-GAN, ARAE, ARIGAN, ArtGAN, ASDL-GAN, ATA-GAN, Attention-GAN, AttGAN, AVID, ...

# Generative Adversarial Network (GAN)



## Variations to GANs

- Several new concepts built on top of GANs have been introduced –
  - InfoGAN – Approximate the data distribution and learn interpretable, useful vector representations of data. 分佈
  - Conditional GANs - Able to generate samples taking into account external information (class label, text, another image). Force G to generate a particular type of output. 解釋

# Contents

- Introduction
- Autoencoder
- Generative Adversarial Networks (GAN)
- Variants of GAN
- Applications of GAN
- Implementation of GAN
- Issues and Conclusions

# Applications of GAN

Labels to Street Scene



input



output

Aerial to Map




input



output

Labels to Facade



input



output

BW to Color



input



output

Day to Night



input



output

Edges to Photo



input



output

# Applications of GAN

- Create characters (Anime)
- Pose Guided Person Image Generation
- Cross-domain transfer
- Clothing creation from images and styles
- Super resolution
- High-resolution image synthesis
- Text to image
- Face synthesis (aging)
- Image inpainting
- Image to image translation (satellite image into map)
- Creating Emoji
- Texture synthesis
- Image (photo) editing
- Video (music) generation


[https://medium.com/@jonathan\\_hui/gan-some-cool-applications-of-gans-4c9ecca35900](https://medium.com/@jonathan_hui/gan-some-cool-applications-of-gans-4c9ecca35900)

# Contents

- Introduction
- Autoencoder
- Generative Adversarial Networks (GAN)
- Variants of GAN
- Applications of GAN
- **Implementation of GAN**
- Issues and Conclusions

# Generative Adversarial Network (GAN)


32






# Generative Adversarial Network (GAN)

## Generator



**New components:**  
Transposed convolution,  
Batch Normalization

## Discriminator



**Binary Classifier:**  
Conv, Leaky ReLU,  
FC, Sigmoid

# Transposed Convolution (Deconvolution)

Conv

Output

Input

Stride=1  
Pad=Valid

Stride=2  
Pad=Valid


Stride=2  
Pad=Same

Deconv

Output

Input

# ReLU vs. Leaky ReLU (LReLU)



Used in Generator

Used in Discriminator

# Batch Normalization (BN)

**Input:** Values of  $x$  over a mini-batch:  $\mathcal{B} = \{x_1 \dots x_m\}$ ;

Parameters to be learned:  $\gamma, \beta$

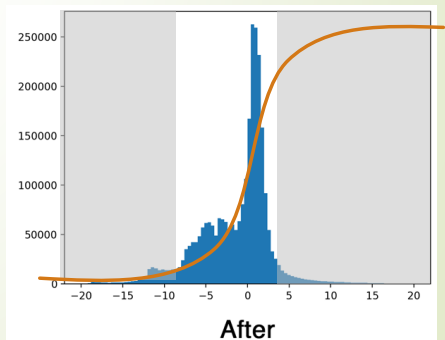

**Output:**  $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$

$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$


$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$

$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$




# Generative Adversarial Network (GAN)

## Generator

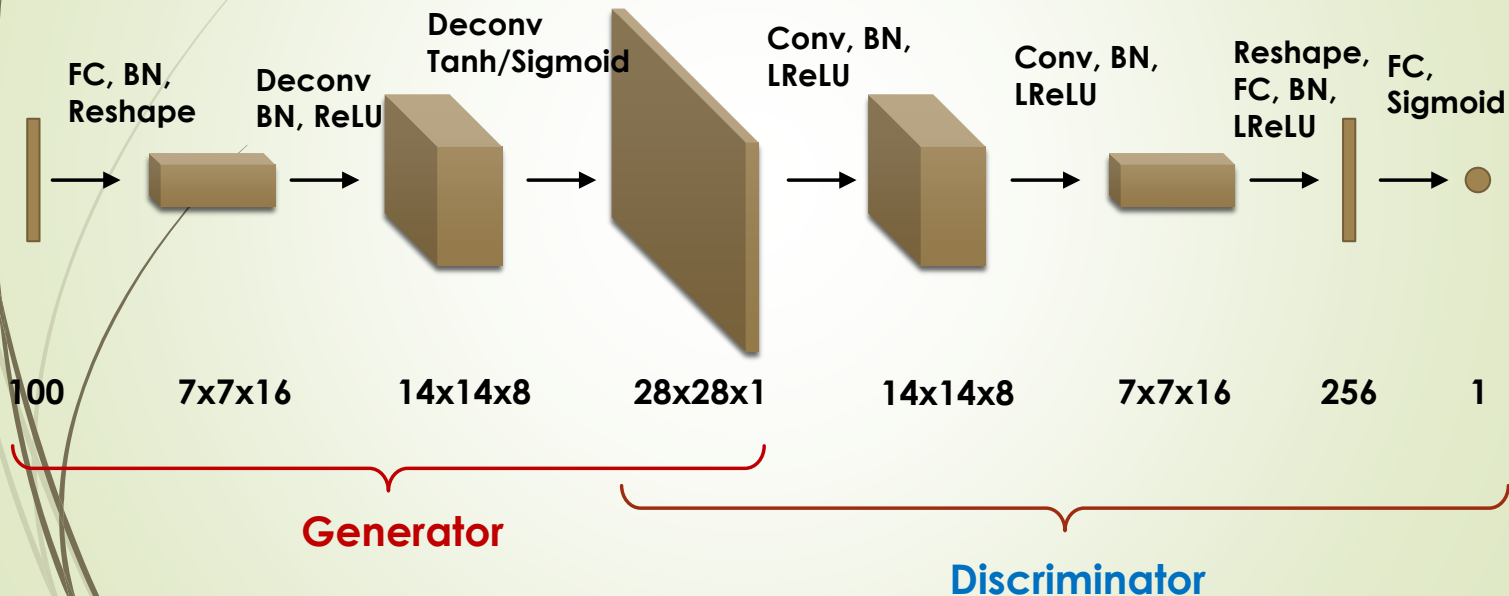


**Deconv (kernel 5x5, stride 2)**  
**BN, ReLU**

## Discriminator

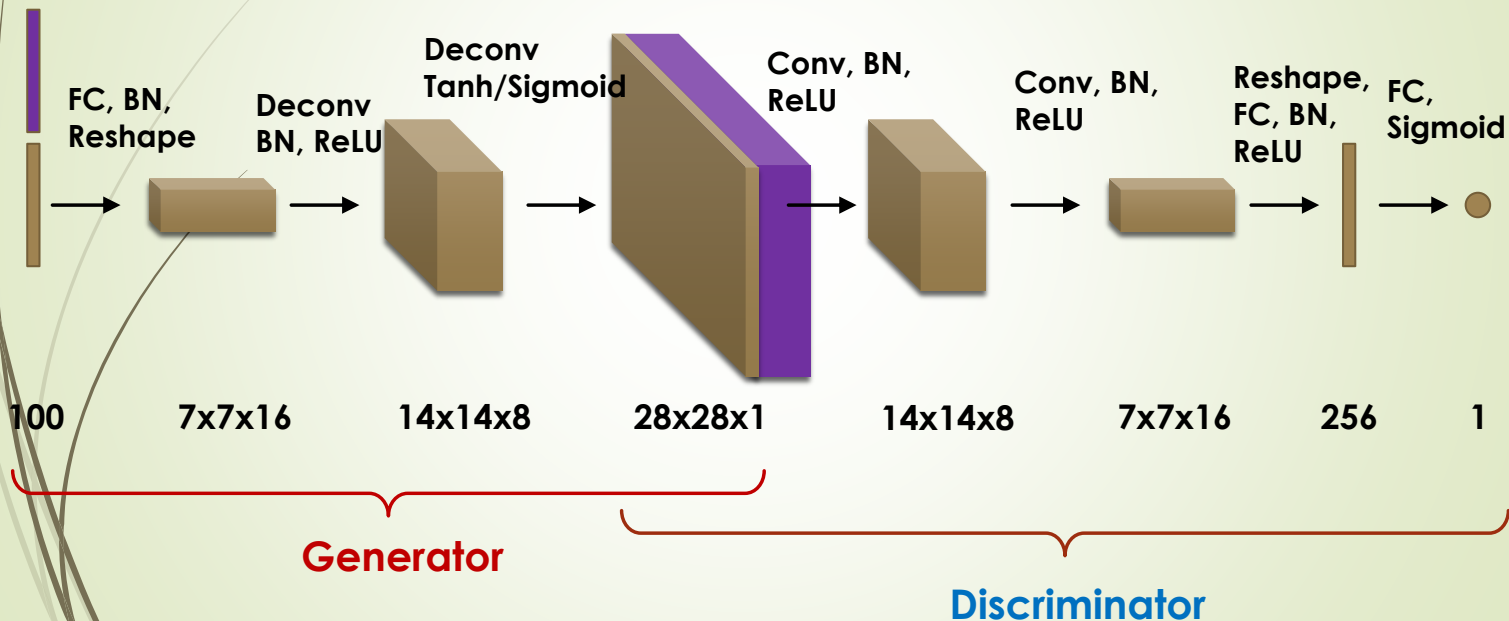


# GAN on MNIST



# Conditional GAN on MNIST

One-hot label



# TensorFlow Implementation of GAN

Run the demo code: (Linux or MacOS)

```
$ git clone https://github.com/carpdm20/DCGAN-tensorflow
```

```
$ cd DCGAN-tensorflow
```

```
$ pip install tqdm (if you do not have this package)
```

```
$ python download.py mnist
```

```
$ python main.py --dataset mnist --input_height=28 --output_height=28 --train
```



# TensorFlow Implementation of GAN

- **Input: z, image, (label)**
- **Network: D, G**
- **Loss: D, G**
- **Optimizer: D, G**

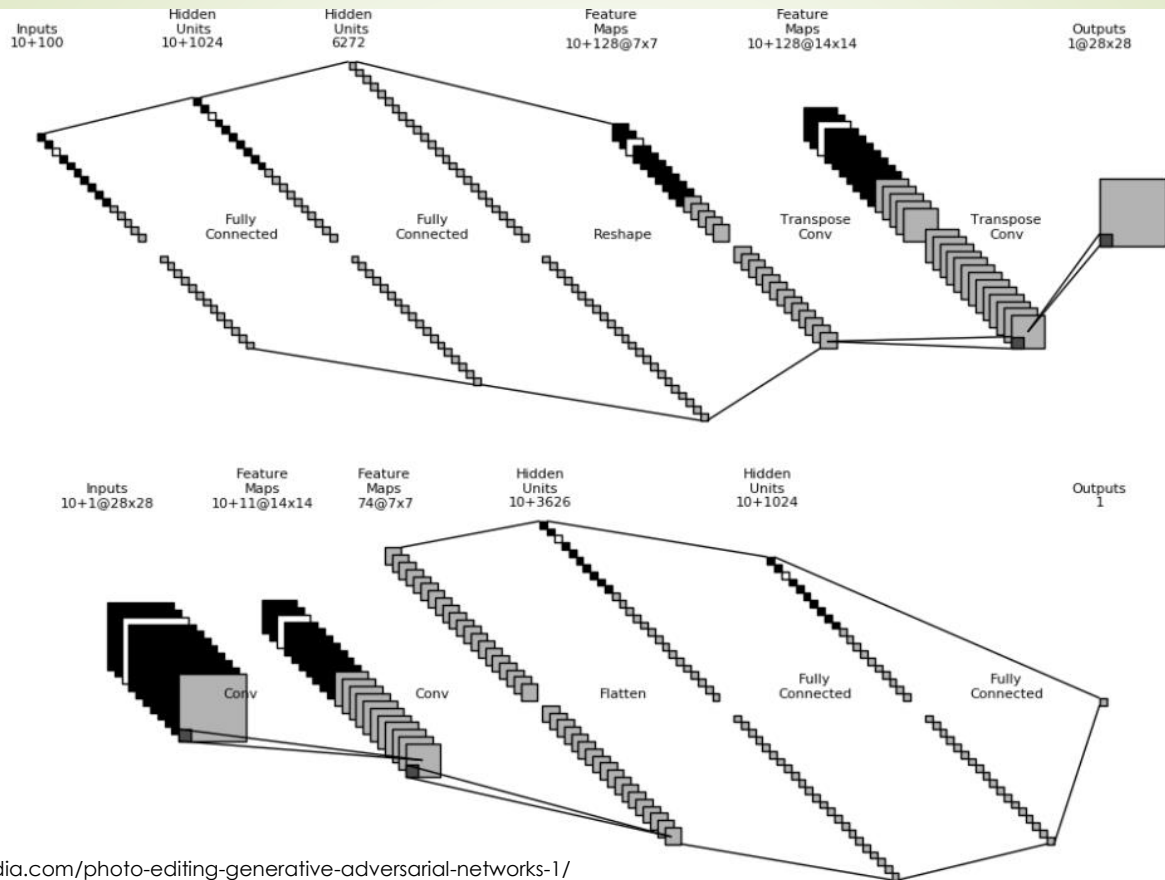
**Training:**

**for epoch**

**for batch**

**Update D**


**Update G**



# TensorFlow Implementation of GAN

## Testing:

- **Random Generation**
- **Conditional Generation**



# Contents

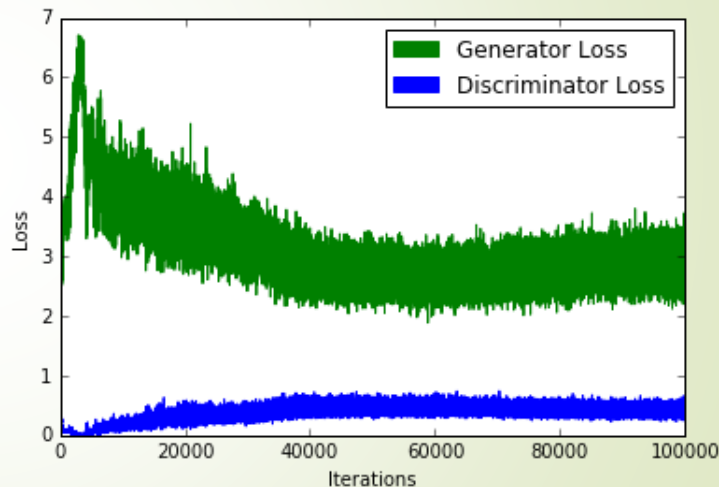
- Introduction
- Autoencoder
- Generative Adversarial Networks (GAN)
- Variants of GAN
- Applications of GAN
- Implementation of GAN
- Issues and Conclusions

## Major Difficulties

- Networks are difficult to converge.
- Ideal goal – Generator and discriminator to reach some desired equilibrium but this is rare.
- GANs are yet to converge on large problems (E.g. Imagenet).

## When to stop training GAN?

- ➔ Generator loss improves when Discriminator loss degrades
- ➔ Discriminator loss improves when Generator loss degrades




## Common Failure Cases

- The discriminator becomes too strong too quickly and the generator ends up not learning anything.
- The generator only learns very specific weaknesses of the discriminator.
- **Mode collapse** in generator
  - The generator learns only a very small subset of the true data distribution.

# Mode Collapse in Generator

- Natural data distributions: multimodal (lots of peaks or modes)
- Generator learns only a limited set of modes to fool the discriminator
- Generator locks on to a few modes and if the discriminator is not fooled, it switches to some other modes, the cycle repeats!





## So what can we do?

- Normalize the inputs
- A modified loss function
- Use a spherical Z
- BatchNorm
- Avoid Sparse Gradients: ReLU, MaxPool
- Use Soft and Noisy Labels
- DCGAN / Hybrid Models
- Track failures early (D loss goes to 0: failure mode)
- If you have labels, use them
- Add noise to inputs, decay over time

# Conclusions

- Train GAN – Use discriminator as base model for transfer learning and the fine-tuning of a production model.
- A well-trained generator has learned the true data distribution well - Use generator as a source of data that is used to train a production model.

## Dive Deeper?

# Generative Adversarial Networks (GANs)

Ian Goodfellow, OpenAI Research Scientist  
NIPS 2016 tutorial  
Barcelona, 2016-12-4

OpenAI

Ian Goodfellow's NIPS 2016 Tutorial

Available online.

# References

- <https://tryolabs.com/blog/2016/12/06/major-advancements-deep-learning-2016/>
- <https://blog.waya.ai/introduction-to-gans-a-boxing-match-b-w-neural-nets-b4e5319cc935#.6l7zh8u50>
- [https://en.wikipedia.org/wiki/Generative\\_adversarial\\_networks](https://en.wikipedia.org/wiki/Generative_adversarial_networks)
- <http://blog.aylien.com/introduction-generative-adversarial-networks-code-tensorflow/>
- <https://github.com/soumith/ganhacks>
- Some cool applications of GAN: [https://medium.com/@jonathan\\_hui/gan-some-cool-applications-of-gans-4c9ecca35900](https://medium.com/@jonathan_hui/gan-some-cool-applications-of-gans-4c9ecca35900)
- Implementation of GAN: <https://docs.google.com/viewer?url=http%3A%2F%2Fweb.eecs.utk.edu%2F~qi%2Fdeeplearning%2Flecture13-gan-implementation.pptx>
- Advances in GAN (including issues and possible solutions) <https://medium.com/beyondminds/advances-in-generative-adversarial-networks-7bad57028032>